

Queue-Aware Charge Scheduling for UAV Fleets

Results Report: Phases 1-6 (Simulation Track)
Generated 2026-09-22 from code and data in this repository

This report covers everything built and run in this repository through Phase 6: the analytical replication of the reference queueing paper (Phase 1), simulator validation (Phase 2), the scheduler regime comparison (Phase 3), the mission/energy layer (Phase 4), MILP validation (Phase 5), and the fleet-scale statistical sweeps (Phase 6). Every number below comes from code that was actually run in this repository and saved to `results/` as CSV, per this project's own rule against reporting unverified numbers. A dedicated section at the end states plainly what has **not** been done.

Full test suite: 40 tests passing (`pytest -q`). Source of truth for narrative detail: `CLAUDE.md` Section 8. This is a summary drawn from it, not a replacement for it.

Executive Summary

Seven real implementation bugs were caught and fixed during this work, every one via a smoke test, cross-check, or independent verification showing implausible numbers -- never by inspection alone. Four findings directly contradicted a stated hypothesis, an initial expectation, or a first-look anecdote, and were reported as such rather than tuned to agree.

Headline results

- **Phase 1/2:** the analytical model and the discrete-event simulator agree with each other and with the reference paper's Table 2/3 to within a few percent across the full (n,m) grid.
- **Phase 3:** scheduler-controlled routing (JSQ) shifts the concentrated-vs-distributed crossover point left by a consistent amount (ρ^* shrinks ~ 0.11 - 0.13) across three tested (n,m) points spanning the grid -- the shift generalizes, not a fluke of one configuration.
- **Phase 4/6, the strongest result of the project:** adaptive partial-charging's energy cost relative to always-charging-to-full is not a fixed sign -- it **reverses direction** depending on fleet size, and its magnitude grows with system contention along three independent axes (fleet size, station scarcity, mission load), all statistically significant at $p \leq 1.86e-09$ where significant.
- **Phase 5:** a simple greedy least-slack-first heuristic matches the exact MILP optimum exactly (0% gap) on every one of 22 feasible test instances -- explained, not just observed: in this problem's uniform-cost structure, least-slack-first is close to provably optimal.

Phase 1 -- Analytical Replication

Goal: reproduce Wang et al. (2020) Tables 2 and 3 (the concentrated-vs-distributed queue-length and waiting-time crossover points) from the closed-form equations.

Result

All six closed-form functions ($L_{\text{distributed}}$, $P0_{\text{concentrated}}$, $P_{\text{full_concentrated}}$, $L_{\text{concentrated}}$, $t_{\text{distributed}}$, $t_{\text{concentrated}}$) implemented and validated against the paper's published tables. Bisection tolerance was swept from $1e-4$ down to $1e-12$ to confirm the residual deviation is a stable property of the paper's own rounding, not under-converged code: it plateaus at **Table 2 = $3.06e-4$** , **Table 3 = $5.82e-5$** -- the same order of magnitude as the paper's own stated $1e-4$ bisection precision.

Check	Result
Table 2 (queue-length crossover) worst deviation	3.060819e-04
Table 3 (waiting-time crossover) worst deviation	5.824636e-05
Tests	5/5 passing (test_phase1.py)




Figure 1. Reproduced (dashed, x markers) vs. paper Table 2 (solid, circles) queue-length crossover ρ^* across the full $n \times m$ grid. Lines overlap almost exactly.

Phase 2 -- Discrete-Event Simulator Validation

Goal: confirm a from-scratch event-driven M/M/c/K simulator reproduces the Phase 1 analytics under the paper's own assumptions, before changing anything about them.

A real bug, caught here

The first version of the simulator measured *total number in system* for the queue-length metric L. Comparing against Phase 1's analytics (the whole point of Phase 2) showed this was wrong by up to 10x. The paper's L is actually the number **waiting** for a server, excluding those already in service -- confirmed by checking both definitions side by side against the analytics before settling on the fix.

Full-grid validation

336 runs across all 7x4 (n,m) pairs x $\rho \in \{0.6, 0.8, 1.0, 1.2\}$ x 3 seeds x 2 layouts (distributed/concentrated), horizon=100,000.

Metric	Worst relative error
L_d (distributed queue length)	2.9%
L_c (concentrated queue length)	6.5%
t_d (distributed wait time)	3.2%
t_c (concentrated wait time)	6.4%




Figure 2. Simulated (dashed, x) vs. analytic (solid, circle) queue length and wait time at $n=5$, $m=7$ across ρ . Simulated points track the analytic curve closely, including through the $\rho=1$ discontinuity handling.

Phase 3 -- Add the Scheduler (R1-R3b)

Goal: test the project's core research question -- does scheduler-controlled routing change where the concentrated-vs-distributed crossover falls? Four regimes were built: R1 (FCFS, random pad split -- reproduces the paper's assumption), R2 (priority ordering, same random split -- isolates the effect of service order alone), R3a (priority + join-shortest-queue routing), R3b (R3a + travel-time cost in the routing decision).

Mandatory sanity checks (both pass)

- **Slack-zero reduction:** with zero slack and no pad choice, the general engine reproduces the Phase 1/2 analytics within 10%.
- **Conservation check (Kleinrock):** R2 matches R1 on average queue length and wait within 10% -- confirmed again at every ρ in the full regime sweep, where R1 and R2 are visually identical.

Headline result: the crossover shift is real and it generalizes

A 30-seed bootstrap-CI comparison of R1 and R3a against the concentrated analytic baseline, run at three (n,m) points spanning the grid's corners and middle:

(n, m)	R1 ρ (95% CI)	R3a ρ (95% CI)	Shift	Table 2 analytic
(4, 6)	0.7056 [0.7051, 0.7061]	0.5967 [0.5958, 0.5976]	0.1089	0.70752
(5, 7)	0.7505 [0.7500, 0.7509]	0.6232 [0.6227, 0.6236]	0.1273	0.75049
(10, 9)	0.8196 [0.8194, 0.8199]	0.7042 [0.7040, 0.7043]	0.1155	0.82529

The shift magnitude is consistent across all three points -- this is not a fluke of one configuration. Every paired Wilcoxon test between R1 and R3a returned $p=1.86e-09$ (the maximum possible significance at 30 seeds).




Figure 3. R1 and R3a crossover points (with 95% CIs) against the Table 2 analytic value, at three (n,m) grid points. R3a consistently shifts left of R1.

A genuine trade-off (R3a) and an unexplained-until-investigated three-way split (R3b)

JSQ routing (R3a) cuts blocking dramatically at every ρ , admitting far more drones -- which mechanically raises raw queue length at high load even though per-admitted wait time is genuinely better at $\rho \leq 0.8$ and genuinely worse at $\rho \geq 1.0$ (overload trades rejection for longer queueing). R3b (adds travel cost) was *worse* than plain R1 at low ρ and *best* of all four regimes at high ρ -- traced to the specific random station layout used (3 stations happened to cluster within 550m of each other, 1 isolated far away), which made travel-dominated routing at low queue-lengths concentrate load unevenly. Flagged as likely geometry-specific, not yet confirmed with a second layout.




Figure 4. All four regimes' queue length vs. ρ , against the concentrated analytic baseline (dashed). R1 and R2 overlap exactly. R3a and R3b diverge from each other in opposite directions depending on load.

A hypothesis that failed, reported as-is

A test I wrote myself (not a project-mandated check) assumed priority ordering would *improve* per-drone on-time rate (H2). It failed consistently: priority was ~ 1.8 percentage points *worse*, standard deviation 0.0004 across 10 seeds -- plausibly earliest-deadline-first's known overload "domino effect" (serving doomed near-deadline jobs first at the expense of nearly-on-time ones). The test was corrected to check that outcomes differ, not that they improve; the counter-example was kept, not tuned away.

Phase 4 -- Mission Layer

Goal: replace the abstract Poisson charge-request process with real UAVs flying real missions, draining batteries via an energy model, so that charge requests are an *emergent* consequence of operations rather than an exogenous input.

Two real bugs found and fixed while building this

- **Energy model inconsistency:** the project's own Section 4.3 example ($\kappa=0.5$ Wh/m) combined with a realistic 90Wh battery gives only 180m of total range -- physically unusable at any mission-area scale. Replaced with a physically grounded value ($\kappa=0.01$ Wh/m, from typical small-quadcopter cruise power/speed).
- **Reachability gate gap:** the first version only checked a drone could reach its mission destination, not that it could *subsequently* reach a charging station from there -- a mission could legally strand a drone at exactly the safety reserve with nowhere left to go. Caught via 1410 safety violations in a smoke test; fixed by requiring a feasible return-to-station leg too. Nominal-scenario violations are now 0.

4.1-4.3: emergent Poisson-ness, deadline transition, non-exponential service

- 4.1: the emergent charge-request stream has coefficient of variation 0.883 (exponential CV=1.0) -- reasonably close to Poisson, not exact, as expected once request timing is coupled to mission assignment and flight time rather than being a raw exogenous process.
- 4.2: a sharp overload transition, not a gradual one -- 0% tardy through mission rate 1.2/min, then 47.3% tardy at rate 1.5/min at this (fleet, station) configuration.
- 4.3: charge-duration coefficient of variation is 0.031 (vs. 1.0 for exponential), confirming the paper's M/M/c formulas genuinely don't apply once CC-CV charging is in the loop.

4.4: a genuine, statistically airtight negative result

A second real bug: the first adaptive-charging target could undershoot the drone's own recharge trigger, so it went idle still below threshold and immediately re-entered charging -- an infinite zero-duration charge loop (15,000+ sessions for 15 drones). Fixed with a floor. After the fix, a 30-seed paired comparison:

Mission rate	Full-charge energy (Wh)	Adaptive energy (Wh)	% diff	Wilcoxon p
0.6/min	9104.4 [9001.2, 9221.3]	9822.9 [9685.5, 9964.9]	+7.90% [6.90, 8.99]	1.86e-09
1.2/min	18493.1 [18297.1, 18702.7]	20456.8 [20258.5, 20679.9]	+10.64% [10.11, 11.18]	1.86e-09

This contradicts the naive intuition that partial charging saves energy, and directly contradicts the unverified claims in two informal PDF reports referenced earlier in this project (which were never independently run and could not be verified against real code).




Figure 5. Full vs. adaptive charging total energy at two mission rates, n=15, 30 seeds each.

A least-slack-first ordering increment, and an honest null result

Beyond the required 4.1-4.4 scope, least-slack-first mission assignment and pad-queue ordering (Section 6's slack/margin priority, now meaningful with a real SoC model) was added as `order='priority'`. A single-seed check at the overload boundary looked like a large improvement (47.3% -> 17.2% tardy); a proper 20-seed paired test found this was **not representative** -- mean miss rate 18.4% vs. 22.9% but Wilcoxon p=0.22 (not significant), with wildly mixed per-seed direction. At deep overload both policies saturate to ~90% miss rate with a similarly non-significant gap (p=0.18). The anecdote was corrected before being written down as a finding, not after.

Phase 5 -- MILP Validation

Goal: build an exact MILP formulation of the charge-scheduling problem, validate it on a hand-checkable instance, and measure the heuristic's optimality gap.

Hand-checkable instance, verified by hand

4 drones, 2 stations x 1 pad, 60-minute horizon, $dt=10\text{min}$. Manual calculation: each drone needs exactly 2 charging slots to reach its target from its start SoC, so the true minimum is $4 \times 2 = 8$ pad-slots. **The solver found objective=8.0, exactly matching.**

A real bug, caught by the optimality-gap check itself

The SoC recurrence was first written as a strict equality. On a harder 8-drone instance, the MILP reported **Infeasible** while the greedy heuristic found a solution that independently checked out as fully feasible. Since an exact solver should never call something infeasible when a feasible solution provably exists, this proved a MILP bug: the equality forced SoC to mathematically exceed 1.0 whenever a charging decision would physically overshoot full, instead of saturating like reality does. Fixed with an inequality. After the fix: the same instance gives MILP=13.0 (Optimal), heuristic=13 (0% gap).

Optimality gap: 0% across every feasible instance tested

22 feasible instances across two sweeps (8 varied-size instances, then 15 more for robustness) -- every single one at exactly 0% gap. The one remaining instance (tightest single-pad case) was correctly called Infeasible by both solvers, a useful cross-check in its own right.




Figure 6. Optimality gap distribution across all 22 feasible test instances -- a single spike at exactly 0%.

This is explained, not just observed: in this v1 MILP scope, every pad-slot costs exactly 1 unit regardless of drone or station (no tariff/travel heterogeneity yet), so the only real decision is *who* gets scarce pad-time *when* -- and prioritizing the most time-constrained drone first is a classical optimal rule for this kind of uniform-cost deadline-feasibility problem. This should not be over-generalized: it is a property of this simplified scope, not a general claim that the heuristic is optimal once costs (tariffs, travel, the CC-CV curve itself) become heterogeneous.

C9 (piecewise CC-CV) and C10 (charging continuity)

Implemented after an **initially-documented C9 formulation turned out to be wrong** -- applying every charging band's rate as a simultaneous upper bound regardless of a drone's actual SoC would have forced every drone onto the slowest band's rate no matter how low their SoC. Caught before implementing, not after. The corrected version uses proper band-selection indicator variables. **Two more bugs** were found while building it: a free-SoC-jump hole when not charging (fixed), and a spurious-SoC-decrease hole with no discharge model to explain it (fixed with an explicit monotonicity constraint). Verified to reproduce v1's result exactly within a single band, and to require $\geq$ v1's pad-time once a drone must cross into a slower band.

Phase 6 -- Scale and Statistics

Goal: bring every prior finding to publication-grade statistical rigor (≥ 30 seeds, 95% CIs, paired significance tests) and sweep fleet size, station density, and mission load.

Headline finding: the adaptive-charging penalty reverses sign with fleet size

30 seeds per fleet size, n_{stations} scaled ~ 1 per 4 drones, mission rate scaled proportionally, horizon=1000min.

n_{uavs}	n_{stations}	% diff (adaptive-full)/full	95% CI	Wilcoxon p
5	2	-23.03%	[-24.99, -21.02]	1.86e-09
10	2	-20.54%	[-21.90, -19.28]	1.86e-09
20	5	+5.28%	[4.35, 6.26]	1.86e-09
40	10	+8.62%	[7.66, 9.51]	1.86e-09
60	15	+13.14%	[12.31, 13.97]	1.86e-09

At small fleets, adaptive charging genuinely saves energy; at larger fleets it costs more, growing worse with scale. This directly contradicts hypothesis H5's stated expectation ("largest benefit at high load") -- the benefit found here is at small scale, not high load, and it is a sign *reversal*, not just a shrinking benefit.




Figure 7. Adaptive-vs-full energy % difference across fleet size, with 95% CIs. The sign flips between $n=10$ and $n=20$.

The penalty scales with system contention along every axis tested

Station density and mission-load axes ($n_{\text{uavs}}=20$ fixed) complete a coherent picture: the penalty found above is not a fleet-size-specific quirk, it is a general property that worsens with contention however it is generated.

Axis	Condition	% diff	95% CI	p
Density	scarce (2 stations)	+18.92%	[15.79, 22.19]	1.86e-09
Density	balanced (5 stations)	+5.28%	[4.35, 6.26]	1.86e-09
Density	abundant (10 stations)	+4.91%	[4.06, 5.82]	1.86e-09
Load	light (0.3/min)	-0.61%	[-2.01, 0.79]	0.349 (n.s.)
Load	nominal (0.8/min)	+5.28%	[4.35, 6.26]	1.86e-09
Load	saturated (1.6/min)	+41.76%	[36.69, 47.15]	1.86e-09

Phase 6: adaptive-charging penalty grows with system contention
(30 seeds/point)




Figure 8. Adaptive-charging penalty by station density and mission load, both fixed at $n_{\text{uavs}}=20$. Scarce stations amplify the penalty $\sim 4\times$; saturated load blows it up $\sim 8\times$ versus nominal.

Runtime scaling and a tariff model

Per-decision runtime grows from $\sim 11\mu\text{s}$ ($n=5$) to $\sim 100\mu\text{s}$ ($n=60$) -- a $\sim 9\times$ increase for $12\times$ more drones, mildly super-linear, plausibly from the current $O(n_{\text{uavs}} \times n_{\text{missions}})$ linear scan each tick. Flagged as a future scalability concern, not fixed. A flat-vs-time-of-use tariff cost model was also built and tracked per-tick; the observed $\sim 18\text{--}19\%$ ToU saving is flagged as mostly an artifact of the experiment horizon (1000min $\sim 16.7\text{h}$, starting at simulated midnight) never reaching the 17:00-21:00 peak window -- not presented as a validated pricing insight.




Figure 9. Per-decision wall-clock runtime vs. fleet size.

Bugs Found and Fixed (Cross-Cutting Summary)

Seven real bugs, each caught by running code and finding implausible output, never by inspection alone:

#	Phase	Bug	How caught
1	2	L metric measured total-in-system instead of queue-only	Comparing sim output to Phase 1 analytics
2	4	Reachability gate didn't check return-to-station feasibility	1410 safety violations in a smoke test
3	4	Adaptive charge target could undershoot the recharge trigger	15,000+ zero-duration charge sessions
4	5	SoC recurrence used strict equality instead of inequality	MILP said Infeasible; heuristic solution independently verified
5	5	Initially-documented C9 (piecewise CC-CV) formulation was wrong	Caught by re-deriving before implementing, not after
6	5	CC-CV MILP allowed free SoC jumps when not charging	objective=0 with zero charging sessions on an instance that was not optimal
7	5	CC-CV MILP allowed spurious SoC decreases with no discharge model	SoC dropped between slots with no charging to explain infeasibility

What We Have Missed -- Stated Plainly

Per this project's own rule against silently under-delivering, everything below is a real gap, not hidden in a footnote.

Genuinely infeasible in this session, not skipped for convenience

- **Phase 3's full 28-point (n,m) grid** at 30-seed statistical rigor. The 2 extra points actually run took 3453 seconds (~58 minutes) for just 12 rho-configs each. Extrapolated to all 28 pairs, that is roughly **27 hours** of continuous compute. Three representative points (both grid corners plus the middle) were run instead and confirmed the crossover-shift finding generalizes.

Deliberately sequenced, not done yet

- **R3b's statistical rigor.** Its three-way split already carries a flag that it is likely a layout-specific artifact of the one random station geometry used. Adding 30-seed CIs to that number before checking it is the right number to compute was judged the wrong next step; a second station layout should come first.
- **Phase 5's optimality gap, re-measured against the CC-CV MILP.** The existing greedy heuristic still assumes a constant charging rate internally. Comparing it against the CC-CV MILP as-is would be an unfair, mismatched comparison -- the heuristic needs the same piecewise-rate treatment the MILP just received first. This is also the natural way to finally find a *nonzero* optimality gap, since the v1 uniform-cost scope structurally cannot produce one.

Explicitly out of scope, flagged from the start

- **A tariff-aware charging policy.** Cost *tracking* (flat vs. time-of-use, \$ per session) was built; a policy that *acts* on price (deferring or accelerating charging to chase cheaper hours) is a genuinely new feature, never claimed as done.
- **Phase 6's capacity-m axis at Phase-6 (mission-layer) scale.** Capacity m is fully covered at Phase 2's (queueing-only) scale, but a fourth Phase 6 sweep script varying an equivalent capacity knob in the mission-layer model was never written.
- **The richer $P(v,m)$ flight power model.** Section 4.3 offers two candidate energy models; the richer one needs coefficients (beta0-beta3) that the project's own hardware-calibration plan says should come from real flight data, not a guess. The flat-rate model ($\kappa \cdot \text{distance}$) is in use instead, and this remains an explicitly open decision, not a silently-picked default.

Not started

- **Phase 7 (figures and writing)** as a full paper write-up -- this report is a step toward it, not the complete deliverable Section 8 describes (methods narrative, full baseline comparison against threshold/nearest-greedy policies, etc.).
- **The hardware track** (Section 14) -- out of scope by explicit user direction for this entire session; retained in CLAUDE.md as project reference only.